This site SFU.ca

CourSys / CMPT 732 G1 / Pages / Assign6

Logged in as mjk30. Logout

## **Assignment 6**

Standalone Reddit Averages

Use the SparkSkeleton and the example "Spark SQL + DataFrames" code from last week to create a program reddit average df.py that calculates the average score for each subreddit (as we have before), but using Spark DataFrames.

We will make a few additions to the code from last week...

When reading JSON into a DataFrame, specifying a schema makes reading much faster: without a schema, the files are read to infer the schema and then again to load the data. Specify the schema when reading the file. We really want to avoid that overhead. Big hint: RedditSchema.

We will write our output as CSV files. The output format won't be exactly like last week, but it will be similar. The code will be like: averages.write.csv(output, mode='overwrite')

**Execution Plans** 

DataFrames have execution plans that describe how they will be computed. Have a look at the plan for this problem (assuming the DataFrame you write is averages):

averages.explain()

Some things to note from the output: when the JSON input is read, you can see which fields are loaded into memory; the average is calculated in three steps. [?]

#### Race!

We now have several implementations for the Reddit averages problem that we could compare: MapReduce, Spark RDDs, and Spark DataFrames.

For each of the Spark implementations, we could compare the usual CPython with PyPy and see if PyPy speeds up the code.

Do these really differ, or did we just waste our time reimplementing the same logic? Feel free to do these runs in small groups: there's not much point to all of us grinding away at the cluster. Write up your answers to the questions separately.

The reddit-6 data set is a yet-larger subset of the Reddit data, and large enough to take at least some measurable time on the cluster. Again, we will limit ourselves to 8 executors to get consistent and fair comparisons.

These commands should get everything to run as described:

export HADOOP\_CLASSPATH=./json-20180813.jar

# Spark DataFrames (with CPython)

# MapReduce time yarn jar al.jar RedditAverage -libjars json-20180813.jar,/opt/hadoop/share/hadoop/tools/lib/lz4-java-1.7.1

time spark-submit --conf spark.dynamicAllocation.enabled=false --num-executors=8 reddit\_average\_df.py /courses/ # Spark RDDs (with CPython) time spark-submit --conf spark.dynamicAllocation.enabled=false --num-executors=8 reddit averages.py /courses/73

# Spark DataFrames (with PyPy) module load spark-pypy time spark-submit --conf spark.dynamicAllocation.enabled=false --num-executors=8 reddit average df.py /courses/

# Spark RDDs (with PyPy) time spark-submit --conf spark.dynamicAllocation.enabled=false --num-executors=8 reddit averages.py /courses/73

For MapReduce, there's nothing on the command line that can limit the job to a certain number of cores, but the input is partitioned into 8 files, so we'll

If you want to see how important specifying a schema was above, you could do another run with the DataFrames implementation without the

schema=... when reading the file. [For me, it was about 60 seconds extra.]

get 8 mappers and one reducer. That should be reasonably comparable to 8 executors and one driver for Spark.

Have a look at the questions below: can you explain what you just observed? [?]

Most-Viewed Wikipedia Pages

As we did with MapReduce and RDDs, let's find the most popular page by hour from the Wikipedia page view statistics.

This time, we will work with the original data format: the hour is **not** in each line of the file, but must be inferred from the filename. The files are named like pagecounts-20160801-120000.gz and the format is like:

en Aaaah 20 231818 en Aaadonta 2 24149 en AaagHiAag 1 8979

These files are in the pagecounts-\* data sets, not the pagecounts-with-time-\* sets.

Use Spark DataFrames to find the the most-viewed page each hour and how many times it was viewed. The program should be called wikipedia\_popular\_df.py and (as usual) take input and output directories from the command line.

As before, we're interested only in: (1) English Wikipedia pages (i.e. language is "en") only. (2) Not the page titled 'Main\_Page'. (3) Not the titles starting with 'Special:'). [The smaller data sets with subsets of pages might not have the main page or special pages: that doesn't mean you aren't responsible for this behaviour.]

You can assume the filenames will be in the format pagecounts-YYYYMMDD-HHMMSS\* (maybe with an extension). We want the YYYYMMDD-HH substring of that as a label for the day/hour.

#### Some hints:

Note the tie in hour 6.

• The DataFrame.read.csv function can read these space-delimited files with a sep argument.

• When reading the file, give a schema argument so data is read correctly with meaningful column names.

• The filename isn't obviously-visible, but can be retrieved with the input file name function if you use it right away: spark.read.csv(...).withColumn('filename', functions.input file name())

• Unless you're more clever than me, converting an arbitrarily-deep filename (that might be file:/// or hdfs:/// and .gz or .lz4 or neither, or who-

knows what else) is very hard in the DataFrame functions. I suggest writing a Python function that takes pathnames and returns string like we need, and using it as a UDF. (Note below.) • To find the most-frequently-accessed page, you'll first need to find the largest number of page views in each hour.

• ... then join that back to the collection of all page counts, so you keep only those with the count == max (count) for that hour. (That means if there's a tie you'll keep both, which is reasonable.)

• As always, use .cache() appropriately. Sort your results by date/hour (and page name if there's a tie) and output as newline-delimited JSON. Part of my output on pagecounts-1 looks like this.

{"hour": "20160801-03", "title": "Simon Pegg", "views": 175} {"hour": "20160801-04", "title": "Simon Pegg", "views": 135} {"hour": "20160801-05", "title": "Simon Pegg", "views": 109} {"hour": "20160801-06", "title": "Simon Cowell", "views": 96} {"hour": "20160801-06", "title": "Simon Pegg", "views": 96} {"hour": "20160801-07", "title": "Simon Pegg", "views": 101}

Your UDF should split the path string on '/', take the last path component, and extract the string we want, like '20160801-12'. You can assume the filename (last path component) will start with 'pagecounts-'.

@functions.udf(returnType=types.StringType()) def path\_to\_hour(path):

with --conf spark.sql.autoBroadcastJoinThreshold=-1.[?]

# **Broadcast Joins & DataFrames**

The join operation that's necessary in this question combines every page and its count (big: there are many Wikipedia pages) with the max views for every hour (small: Wikipedia has only existed for a few thousand days × 24 hours). That seems like a candidate for the broadcast join we saw in the last exercise.

With DataFrames, we can get that optimization easily: either by the runtime noticing that one of the join inputs is small and automatically broadcasting, or by explicitly flagging a DataFrame as broadcast-able.

Try your code on a reasonably-large input set with and without the broadcast hint. Use an input set that's large enough to see what would happen on big data. That might be pagecounts - 3 on the cluster. [?]

spark.sql.autoBroadcastJoinThreshold=-1 Also, compare the execution plans for your final result (from .explain()) for the join (1) with the broadcast hint and (2) without the broadcast hint and

There is an automatic detection of broadcast-able DataFrames, which you may have to disable on the command line to see the difference: --conf

Weather and Temperature Ranges 1: DataFrame Methods

Let's look again at the Global Historical Climatology Network data. Their archives contain (among other things) lists of weather observations from various weather stations formatted as CSV files like this:

US1FLSL0019,20130101,SNOW,0,,,N, US1TXTV0133,20130101,PRCP,30,,,N, USC00178998,20130101,TMAX,-22,,,7,0700 USC00178998,20130101,TMIN,-117,,,7,0700 USC00178998,20130101,TOBS,-28,,,7,0700 The fields are the weather station; the date (Jan 1 2013 in these lines); the observation (min/max temperature, amount of precipitation, etc); the value (an

US1FLSL0019,20130101,PRCP,0,,,N,

integer); and several flags about the observation. Their readme file explains the fields in detail. We will worry about the min and max temperature (TMIN, TMAX) which are given as °C × 10. As we did last week we only want to keep the values where

the QFLAG field is null, indicating a "correct" observation. Data sets are weather-1 (and -2 and -3), which are subsets of the full data, partitioned nicely. Create a program temp range.py and take input and

The CSV files should be read by using spark.read.csv() and specifying a schema, as before.

## The Problem What I want to know is: what weather station had the largest temperature difference on each day? That is, where was the largest difference

output directory arguments, as usual.

between TMAX and TMIN? We want final data like this (if we .show(10)):

+----+ date station|range|

+----+ |20170501|USC00410779| 26.1| |20170502|USC00413411| 26.1| |20170503|RSM00031939| 24.6| |20170504|USS0013C39S| 22.5| |20170505|USW00094012| 27.8| |20170506|USC00243110| 31.1| |20170507|RSM00030521| 27.1| |20170508|RSM00030971| 25.2| |20170509|USC00047109| 23.9| |20170509|USW00094107| 23.9| +----+ Output

### The output should be sorted by date and by station if there is a tie. Our final output in this case should be the same as the input: CSV files in the default way that Spark produces.

**Getting There** For this question, I'm going to ask you to implement this with the **Python API** (that is, with methods on DataFrames, not spark.sql()), but the next

by 10 so we have actual °C values. Find the max. Then join the max values back to find out which station(s) that range came from. If there's a tie, this method will get us both results, which seems reasonable. There are a few points where you should be caching DataFrames here: make sure you do.

First get the TMIN and TMAX for each day separately, and then join them back together and subtract to get the temperature range. At some point, divide

It's possible that doing a broadcast join here would be faster than the standard join: you can include the broadcast hint if you want, but don't worry too much about it.

part is to do the same thing with the SQL syntax: start with whichever you like, obviously.

Weather and Temperature Ranges 2: SQL Syntax Create a temp\_range\_sql.py that does the same as the above, but using the SQL syntax. That is, most of the work should be done with calls to

spark.sql(), not by calling methods on DataFrame objects. The logic should translate more-or-less directly. You just have to spell it with Spark's SQL dialect. [?]

The output should be exactly the same as the DataFrames-methods implementation.

Questions

# In a text file answers.txt, answer these questions:

1. In the Reddit averages execution plan, which fields were loaded? How was the average computed (and was a combiner-like step done)?

Submit your files to the CourSys activity Assignment 6.

2. What was the running time for your Reddit averages implementations in the five scenarios described above? How much difference did Python implementation make (PyPy vs the default CPython)? Why was it large for RDDs but not for DataFrames? 3. How much of a difference did the broadcast hint make to the Wikipedia popular code's running time (and on what data set)?

4. How did the Wikipedia popular execution plan differ with and without the broadcast hint? 5. For the weather data question, did you prefer writing the "DataFrames + Python methods" style, or the "temp tables + SQL syntax" style form solving

the problem? Which do you think produces more readable code?

Submission

Terms and conditions

© Simon Fraser University

Community

About

Give to SFU

Canada V5A 1S6